

```
In [ ]: import tensorflow as tf
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import StandardScaler
        from sklearn import metrics
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
        %matplotlib inline
        from tqdm.notebook import tqdm
        import warnings
        warnings.filterwarnings("ignore")
```

```
In [ ]: boston = tf.keras.datasets.boston_housing
```

```
In [ ]: dir(boston)
```

```
Out[ ]: ['__builtins__',
        '__cached__',
        '__doc__',
        '__file__',
        '__loader__',
        '__name__',
        '__package__',
        '__path__',
        '__spec__',
        'load_data']
```

```
In [ ]: boston_data = boston.load_data()
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston\\_housing.npz](https://storage.googleapis.com/tensorflow/tf-keras-datasets/boston_housing.npz)  
 57026/57026 [=====] - 0s 0us/step

```
In [ ]: (x_train, y_train), (x_test, y_test) = tf.keras.datasets.boston_housing.load_data(p
```

```
In [ ]: x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[ ]: ((404, 13), (404,)), (102, 13), (102,))
```

```
In [ ]: scaler = StandardScaler()
```

```
In [ ]: x_train_scaled = scaler.fit_transform(x_train)
        x_test_scaled = scaler.transform(x_test)

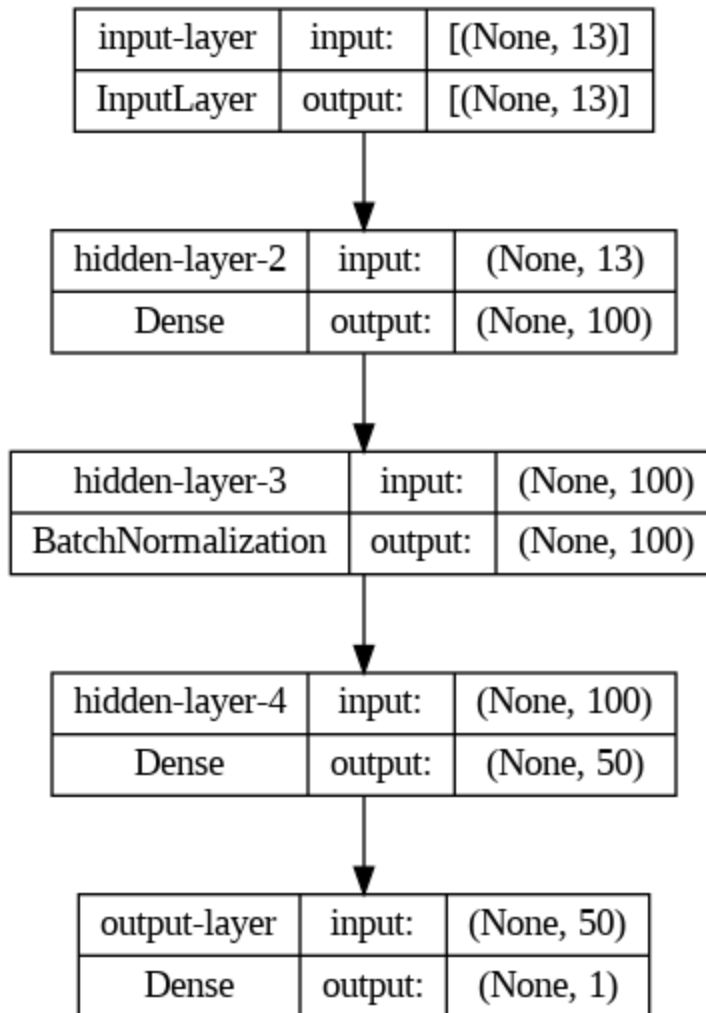
        y_train_scaled = scaler.fit_transform(y_train.reshape(-1, 1))
        y_test_scaled = scaler.transform(y_test.reshape(-1, 1))
```

```
In [ ]: model = tf.keras.models.Sequential([
        tf.keras.layers.Input(shape=(13), name='input-layer'),
        tf.keras.layers.Dense(100, name='hidden-layer-2'),
        tf.keras.layers.BatchNormalization(name='hidden-layer-3'),
        tf.keras.layers.Dense(50, name='hidden-layer-4'),
```

```
tf.keras.layers.Dense(1, name='output-layer')
])
```

```
In [ ]: tf.keras.utils.plot_model(model, show_shapes=True)
```

```
Out[ ]:
```



```
In [ ]: model.summary()
```

Model: "sequential\_10"

Layer (type)	Output Shape	Param #
hidden-layer-2 (Dense)	(None, 100)	1400
hidden-layer-3 (BatchNormalization)	(None, 100)	400
hidden-layer-4 (Dense)	(None, 50)	5050
output-layer (Dense)	(None, 1)	51

=====  
Total params: 6901 (26.96 KB)  
Trainable params: 6701 (26.18 KB)  
Non-trainable params: 200 (800.00 Byte)

```
In [ ]: model.compile(  
    optimizer='adam',  
    loss='mse',  
    metrics=['mae']  
)
```

```
In [ ]: history = model.fit(x_train, y_train, batch_size=32, epochs=20, validation_data=(x_
```

```
Epoch 1/20
13/13 [=====] - 1s 18ms/step - loss: 574.8789 - mae: 22.363
4 - val_loss: 2224.5266 - val_mae: 44.4160
Epoch 2/20
13/13 [=====] - 0s 7ms/step - loss: 550.9838 - mae: 22.0251
- val_loss: 705.3503 - val_mae: 24.4563
Epoch 3/20
13/13 [=====] - 0s 7ms/step - loss: 530.6443 - mae: 21.6095
- val_loss: 1005.8787 - val_mae: 30.3659
Epoch 4/20
13/13 [=====] - 0s 6ms/step - loss: 501.7632 - mae: 21.0424
- val_loss: 629.7626 - val_mae: 23.7866
Epoch 5/20
13/13 [=====] - 0s 7ms/step - loss: 462.1404 - mae: 20.1391
- val_loss: 658.0805 - val_mae: 24.4096
Epoch 6/20
13/13 [=====] - 0s 6ms/step - loss: 399.1111 - mae: 18.7222
- val_loss: 542.7188 - val_mae: 22.2311
Epoch 7/20
13/13 [=====] - 0s 7ms/step - loss: 312.8028 - mae: 16.4244
- val_loss: 313.3749 - val_mae: 16.6600
Epoch 8/20
13/13 [=====] - 0s 6ms/step - loss: 209.3056 - mae: 13.1072
- val_loss: 347.4641 - val_mae: 17.6651
Epoch 9/20
13/13 [=====] - 0s 7ms/step - loss: 112.0523 - mae: 8.9190
- val_loss: 192.2844 - val_mae: 12.5691
Epoch 10/20
13/13 [=====] - 0s 7ms/step - loss: 55.1853 - mae: 5.6105 -
val_loss: 115.9377 - val_mae: 9.1478
Epoch 11/20
13/13 [=====] - 0s 6ms/step - loss: 33.7891 - mae: 4.2850 -
val_loss: 29.2040 - val_mae: 3.8567
Epoch 12/20
13/13 [=====] - 0s 7ms/step - loss: 36.9304 - mae: 4.4974 -
val_loss: 99.4068 - val_mae: 8.0226
Epoch 13/20
13/13 [=====] - 0s 6ms/step - loss: 32.5308 - mae: 4.2865 -
val_loss: 35.6296 - val_mae: 4.8049
Epoch 14/20
13/13 [=====] - 0s 6ms/step - loss: 29.5778 - mae: 4.0588 -
val_loss: 41.9396 - val_mae: 4.5167
Epoch 15/20
13/13 [=====] - 0s 7ms/step - loss: 29.4901 - mae: 4.0193 -
val_loss: 49.0794 - val_mae: 5.6884
Epoch 16/20
13/13 [=====] - 0s 6ms/step - loss: 32.3794 - mae: 4.1094 -
val_loss: 26.8985 - val_mae: 3.5455
Epoch 17/20
13/13 [=====] - 0s 7ms/step - loss: 31.6614 - mae: 4.0583 -
val_loss: 33.6451 - val_mae: 3.8305
Epoch 18/20
13/13 [=====] - 0s 7ms/step - loss: 26.4994 - mae: 3.7155 -
val_loss: 26.4998 - val_mae: 3.4068
Epoch 19/20
13/13 [=====] - 0s 7ms/step - loss: 27.2167 - mae: 3.8110 -
```


val\_loss: 29.5473 - val\_mae: 4.1069

Epoch 20/20

13/13 [=====] - 0s 7ms/step - loss: 27.8241 - mae: 3.9048 -

val\_loss: 27.6318 - val\_mae: 3.7487


```
In [ ]: pd.DataFrame(history.history).plot(figsize=(10,7))
plt.title("Metrics graph")
plt.show()
```



```
In [ ]: y_pred = model.predict(x_test)
```

4/4 [=====] - 0s 4ms/step

```
In [ ]: sns.regplot(x=y_test, y=y_pred)
plt.title("Regression Line for Predicted values")
plt.show()
```



```
In [ ]: def regression_metrics_display(y_test, y_pred):  
        print(f"MAE is {metrics.mean_absolute_error(y_test, y_pred)}")  
        print(f"MSE is {metrics.mean_squared_error(y_test, y_pred)}")  
        print(f"R2 score is {metrics.r2_score(y_test, y_pred)}")
```

```
In [ ]: regression_metrics_display(y_test, y_pred)
```

```
MAE is 3.7487237733953136  
MSE is 27.631791791938596  
R2 score is 0.6119240015553091
```