

chit-7-fashion

May 6, 2025

```
[2]: import tensorflow as tf
      #from tensorflow.keras.layers.experimental import preprocessing
      from sklearn.model_selection import train_test_split
      from mlxtend.plotting import plot_confusion_matrix
      from sklearn import metrics
      import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns
      %matplotlib inline
      from tqdm.notebook import tqdm
      import random
      import warnings
      warnings.filterwarnings("ignore")

[3]: (trainX, trainY), (testX, testY) = tf.keras.datasets.fashion_mnist.load_data()
      trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))
      testX = testX.reshape((testX.shape[0], 28, 28, 1))
      trainY_cat = tf.keras.utils.to_categorical(trainY)
      testY_cat = tf.keras.utils.to_categorical(testY)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-labels-idx1-ubyte.gz
29515/29515          0s 9us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-images-idx3-ubyte.gz
26421880/26421880    9s
0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-labels-idx1-ubyte.gz
5148/5148           0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-images-idx3-ubyte.gz
4422102/4422102      3s
1us/step
```

```
[4]: train_norm = trainX.astype('float32')
test_norm = testX.astype('float32')
train_norm = train_norm / 255.0
test_norm = test_norm / 255.0
```

```
[5]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', '␣',
↪ 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
```

```
[6]: plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(trainX[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[trainY[i]])
plt.show()
```



```
[7]: model = tf.keras.models.Sequential([
    # First convolutional layer
    tf.keras.layers.Conv2D(64, kernel_size=(3, 3), input_shape=(28, 28, 1),
    ↪activation='relu', name='conv-layer-1'),

    # First pooling layer
    tf.keras.layers.AvgPool2D(pool_size=(2, 2), name='pooling-layer-1'),

    # Second convolutional layer
    tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu',
    ↪padding='same', name='conv-layer-2'),
```

```

# Second pooling layer
tf.keras.layers.AvgPool2D(pool_size=(2, 2), name='pooling-layer-2'),

# Global average pooling
tf.keras.layers.GlobalAveragePooling2D(name='pooling-layer-3'),

# Fully connected layer with softmax activation for multi-class
classification
tf.keras.layers.Dense(len(class_names), activation="softmax",
name="output-layer")
])

```

```

[8]: model.compile(loss="categorical_crossentropy",
optimizer="adam",
metrics=["accuracy"])

```

```

[9]: history = model.fit(train_norm, trainY_cat, epochs=10,
validation_data=(test_norm, testY_cat))

```

```

Epoch 1/10
1875/1875          13s 7ms/step -
accuracy: 0.4759 - loss: 1.4343 - val_accuracy: 0.7090 - val_loss: 0.8380
Epoch 2/10
1875/1875          13s 7ms/step -
accuracy: 0.7230 - loss: 0.7909 - val_accuracy: 0.7534 - val_loss: 0.7231
Epoch 3/10
1875/1875          13s 7ms/step -
accuracy: 0.7497 - loss: 0.7128 - val_accuracy: 0.7665 - val_loss: 0.6767
Epoch 4/10
1875/1875          14s 8ms/step -
accuracy: 0.7677 - loss: 0.6566 - val_accuracy: 0.7369 - val_loss: 0.7168
Epoch 5/10
1875/1875          15s 8ms/step -
accuracy: 0.7777 - loss: 0.6333 - val_accuracy: 0.7628 - val_loss: 0.6469
Epoch 6/10
1875/1875          16s 9ms/step -
accuracy: 0.7849 - loss: 0.6138 - val_accuracy: 0.7846 - val_loss: 0.6163
Epoch 7/10
1875/1875          17s 9ms/step -
accuracy: 0.7922 - loss: 0.5906 - val_accuracy: 0.7991 - val_loss: 0.5942
Epoch 8/10
1875/1875          16s 9ms/step -
accuracy: 0.7973 - loss: 0.5740 - val_accuracy: 0.7923 - val_loss: 0.5979
Epoch 9/10
1875/1875          16s 9ms/step -
accuracy: 0.8038 - loss: 0.5531 - val_accuracy: 0.8059 - val_loss: 0.5672
Epoch 10/10

```

```
1875/1875          16s 9ms/step -
accuracy: 0.8126 - loss: 0.5369 - val_accuracy: 0.8096 - val_loss: 0.5543
```

```
[10]: tf.keras.utils.plot_model(model, show_shapes=True)
```

You must install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for `plot_model` to work.

```
[11]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
Param #		
conv-layer-1 (Conv2D)	(None, 26, 26, 64)	
↳ 640		
pooling-layer-1 (AveragePooling2D)	(None, 13, 13, 64)	
↳ 0		
conv-layer-2 (Conv2D)	(None, 13, 13, 32)	
↳ 18,464		
pooling-layer-2 (AveragePooling2D)	(None, 6, 6, 32)	
↳ 0		
pooling-layer-3	(None, 32)	
↳ 0		
(GlobalAveragePooling2D)		
↳		
output-layer (Dense)	(None, 10)	
↳ 330		

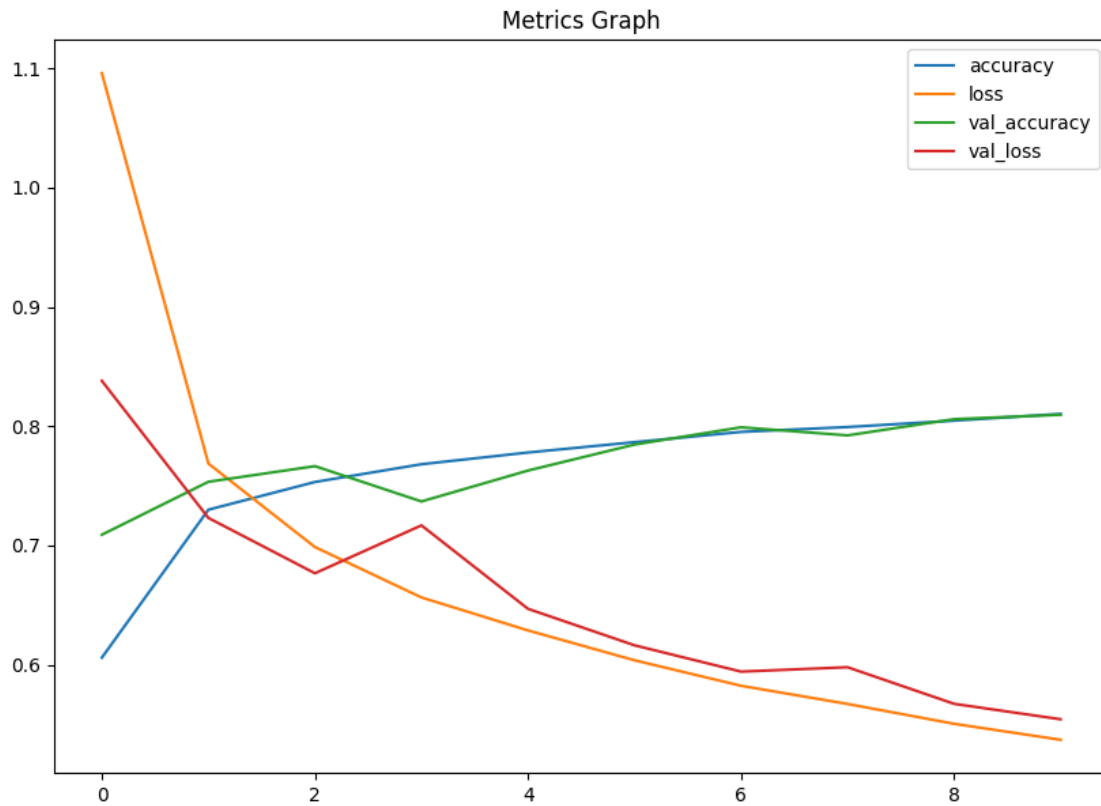
Total params: 58,304 (227.75 KB)

Trainable params: 19,434 (75.91 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 38,870 (151.84 KB)

```
[12]: pd.DataFrame(history.history).plot(figsize=(10,7))
plt.title("Metrics Graph")
plt.show()
```



```
[13]: model.evaluate(testX, testY_cat)
```

```
313/313          1s 3ms/step -
accuracy: 0.6350 - loss: 115.7895
```

```
[13]: [119.34942626953125, 0.6317999958992004]
```

```
[14]: predictions = model.predict(testX)
```

```
313/313          1s 3ms/step
```

```
[15]: predictions = tf.argmax(predictions, axis=1)
```

```
[16]: y_test = tf.argmax(testY_cat, axis=1)
```

```
[17]: y_test = tf.Variable(y_test)
```


```
[18]: print(metrics.accuracy_score(y_test, predictions))
```

0.6318

```
[19]: print(metrics.classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.41	0.58	0.48	1000
1	0.53	0.98	0.69	1000
2	0.71	0.64	0.67	1000
3	0.67	0.15	0.25	1000
4	0.45	0.81	0.57	1000
5	0.95	0.53	0.68	1000
6	0.47	0.24	0.32	1000
7	0.96	0.60	0.74	1000
8	0.92	0.84	0.88	1000
9	0.81	0.95	0.88	1000
accuracy			0.63	10000
macro avg	0.69	0.63	0.61	10000
weighted avg	0.69	0.63	0.61	10000

```
[20]: cm = metrics.confusion_matrix(y_test, predictions)
plot_confusion_matrix(cm, figsize=(10,7), class_names=class_names)
plt.title("Confusion Matrix")
plt.show()
```



```
[21]: images = []
labels = []
random_indices = random.sample(range(len(testX)), 10)
for idx in random_indices:
    images.append(testX[idx])
    labels.append(testY_cat[idx])
images = np.array(images)
labels = np.array(labels)
fig = plt.figure(figsize=(20, 8))
rows = 2
cols = 5
x = 1
```



```

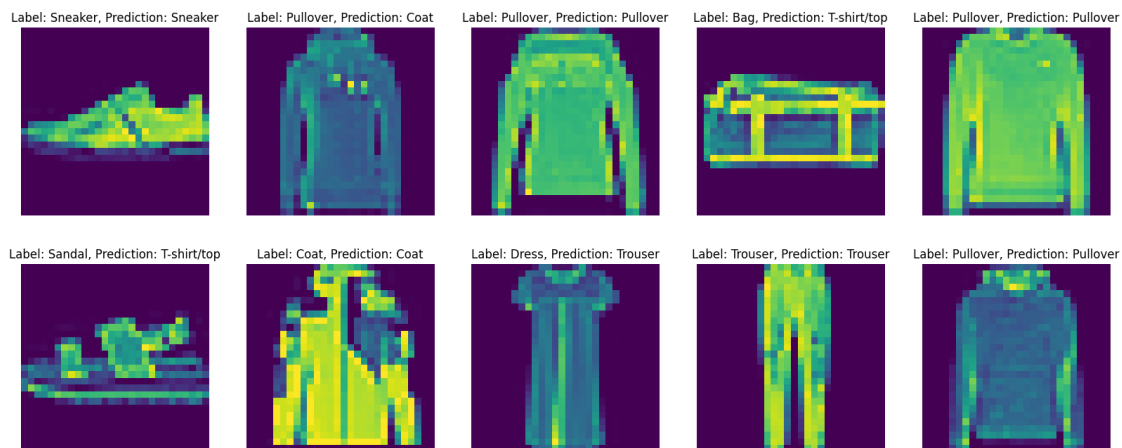
for image, label in zip(images, labels):
    fig.add_subplot(rows, cols, x)
    prediction = model.predict(tf.expand_dims(image, axis=0))
    prediction = class_names[tf.argmax(prediction.flatten())]
    label = class_names[tf.argmax(label)]
    plt.title(f"Label: {label}, Prediction: {prediction}")
    plt.imshow(image/255.)
    plt.axis("off")
    x += 1

```

```

1/1      0s 30ms/step
1/1      0s 26ms/step
1/1      0s 23ms/step
1/1      0s 23ms/step
1/1      0s 20ms/step
1/1      0s 19ms/step
1/1      0s 20ms/step
1/1      0s 19ms/step
1/1      0s 18ms/step
1/1      0s 19ms/step

```



[]: